

# Volatility Forecasting Report

Generated: 2025-11-07 23:28:57

Author: PhD Research Team

---

## Table of Contents

1. HAR-X Model Results
    - Exogenous Variables
    - HARX Model Performance
    - HAR-X Model Predictions vs True RV
    - Loss Metrics Over Time
  2. HAR-X Ensemble Model Results
    - Ensemble Weights
    - HAR-X Ensemble Performance Summary
    - HAR-X Ensemble Predictions and Loss Metrics
    - HAR-X Diebold-Mariano Test
  3. Machine Learning Models
    - RF Model Performance
      - RF Predictions vs True RV
      - RF Loss Metrics Over Time
    - GBM Model Performance
      - GBM Predictions vs True RV
      - GBM Loss Metrics Over Time
    - XGBOOST Model Performance
      - XGBOOST Predictions vs True RV
      - XGBOOST Loss Metrics Over Time
    - LIGHTGBM Model Performance
      - LIGHTGBM Predictions vs True RV
      - LIGHTGBM Loss Metrics Over Time
    - CATBOOST Model Performance
      - CATBOOST Predictions vs True RV
      - CATBOOST Loss Metrics Over Time
  4. Temporal Fusion Transformer (TFT) Results
- 

## Executive Summary

This report presents a comprehensive analysis of volatility forecasting using Heterogeneous Autoregressive (HAR) and HAR with exogenous variables (HAR-X) models applied to Treasury Bond ETF (TLT) data.

**Key Objectives:** - Evaluate multiple volatility estimators (Squared Return, Parkinson, Garman-Klass, Rogers-Satchell) - Compare HAR and HAR-X model

performance across different rolling window sizes - Implement ensemble forecasting using inverse QLIKE weighting - Validate models using out-of-sample testing and statistical comparisons

**Main Findings:** - Ensemble models outperform individual estimators across all metrics - HAR-X with window=756 provides most stable and consistent predictions - Exogenous variables offer marginal but meaningful improvement in forecast calibration - No statistically significant difference between windows 504 and 756 (DM test) - Both models demonstrate strong forecasting ability with well-behaved residuals

## Data Description

**Dataset:** iShares 20+ Year Treasury Bond ETF (TLT)

**Period:** 2003-01-01 to 2024-12-30

**Frequency:** Daily

**Total Observations:** 5536

**Price Data Components:** - Open, High, Low, Close prices - Trading volume  
- Adjusted close prices

**Target Variable:** - Realized Volatility (RV): Annualized variance computed from log returns - Log-transformed for modeling to ensure stationarity

**Train/Test Split:** - Training Set: 2003-01-06 to 2022-12-30 (data before 2023-01-01) — 5033 observations - Test Set: 2023-01-03 to 2024-12-30 (data on/after 2023-01-01) — 501 observations

## Methodology

### Model Framework

#### 1. Volatility Estimation

Four volatility estimators are computed from OHLC data: - **Squared Return (RV):**  $RV^2 = 252 \times (\log(C/O))^2$  - **Parkinson:**  $RV^2 = 252 \times (1/(4\ln 2)) \times (\log(H/L))^2$  - **Garman-Klass:**  $RV^2 = 252 \times [0.5(\log(H/L))^2 - (2\ln 2 - 1)(\log(C/O))^2]$  - **Rogers-Satchell:**  $RV^2 = 252 \times [\log(H/O)\log(H/C) + \log(L/O)\log(L/C)]$

All estimators are log-transformed for modeling.

#### 2. HAR Model

The HAR model captures heterogeneous volatility components:

$$\log(RV) = \alpha + \beta_1 \cdot RV_{\text{Daily}} + \beta_2 \cdot RV_{\text{Weekly}} + \beta_3 \cdot RV_{\text{Monthly}}$$

Where: -  $RV_{\text{Daily}}$  : Daily component (lag 1) -  $RV_{\text{Weekly}}$  : Weekly component (5-day average) -  $RV_{\text{Monthly}}$  : Monthly component (22-day average)

#### 3. HAR-X Model

$$\log(\text{RV}) = \beta_0 + \beta_1 \cdot \text{RV} + \beta_2 \cdot \text{RV}^2 + \beta_3 \cdot \text{RV}^3 + \beta_4 \cdot \Sigma + \beta_5 \cdot X + \epsilon$$

## 4. Rolling Window Estimation

## 5. Ensemble Forecasting

$$w = (1/Q_{\text{LIKE}}) / \Sigma (1/Q_{\text{LIKE}})$$

## 6. Evaluation Metrics

- ## Volatility Estimators Analysis

The following table presents the stationarity and autocorrelation diagnostics for each volatility estimator. We use the Augmented Dickey-Fuller (ADF) test and Kwiatkowski-Phillips-Schmidt-Shin (KPSS) test to assess stationarity, and the Ljung-Box test to check for serial correlation in the residuals.

| Estimator         | ADF stat | ADF p | ADF pass (p ) | KPSS stat |
|-------------------|----------|-------|---------------|-----------|
| square_est_log    | -9.175   | 0     | True          | 1.852     |
| parkinson_est_log | -5.068   | 0     | True          | 0.876     |
| gk_est_log        | -4.993   | 0     | True          | 0.927     |
| rs_est_log        | -5.862   | 0     | True          | 1.251     |

| Estimator      | KPSS p | KPSS pass (p> ) | LB p @10 | LB p @20 |
|----------------|--------|-----------------|----------|----------|
| square_est_log | 0.01   | False           | 0        | 0        |

| Estimator         | KPSS p | KPSS pass (p> ) | LB p @10 | LB p @20 |
|-------------------|--------|-----------------|----------|----------|
| parkinson_est_log | 0.01   | False           | 0        | 0        |
| gk_est_log        | 0.01   | False           | 0        | 0        |
| rs_est_log        | 0.01   | False           | 0        | 0        |

**Table 1: Diagnostic Tests for Volatility Estimators (Part 3/3)**

| Estimator         | White noise (LB) | Stationary (ADF KPSS) |
|-------------------|------------------|-----------------------|
| square_est_log    | 0                | 0                     |
| parkinson_est_log | 0                | 0                     |
| gk_est_log        | 0                | 0                     |
| rs_est_log        | 0                | 0                     |

### ACF and PACF Analysis

The Autocorrelation Function (ACF) and Partial Autocorrelation Function (PACF) plots reveal: - **Slow decay in ACF**: Indicates long memory in volatility, consistent with volatility clustering - **Significant PACF spikes**: Suggests short-term AR effects up to 5-15 lags - **HAR model justification**: These patterns support using daily (1), weekly (5), and monthly (22) lags

### ACF and PACF for All Volatility Estimators

## HAR Model Results

### Performance Metrics Across Windows

The HAR model was evaluated across multiple rolling window sizes: 252, 504, 756, 1008, and 1260 days. Below are the comprehensive performance metrics for each estimator and window size.

**Table 2: HAR Model Performance Summary (QLIKE and MSPE)**

|                            | QLIKE_mean | QLIKE_std | MSPE_mean   | MSPE_std    |
|----------------------------|------------|-----------|-------------|-------------|
| (252, 'square_est_log')    | -0.769     | 12.111    | 1.41257e+17 | 3.76377e+18 |
| (252, 'parkinson_est_log') | -1.198     | 8.147     | 1.6656e+17  | 4.56201e+18 |
| (252, 'gk_est_log')        | -1.162     | 8.635     | 1.90755e+17 | 6.10036e+18 |
| (252, 'rs_est_log')        | -1.092     | 9.508     | 2.15613e+17 | 7.36021e+18 |
| (504, 'square_est_log')    | -0.91      | 9.818     | 1.30743e+17 | 3.60831e+18 |
| (504, 'parkinson_est_log') | -1.167     | 7.949     | 1.2193e+17  | 2.64853e+18 |
| (504, 'gk_est_log')        | -1.161     | 7.916     | 1.21592e+17 | 2.70329e+18 |
| (504, 'rs_est_log')        | 2.677      | 241.793   | 9.89925e+16 | 1.90839e+18 |



|                             | QLIKE_mean | QLIKE_std | MSPE_mean   | MSPE_std    |
|-----------------------------|------------|-----------|-------------|-------------|
| (756, 'square_est_log')     | -0.712     | 11.015    | 1.20577e+17 | 2.87891e+18 |
| (756, 'parkinson_est_log')  | -1.074     | 8.005     | 1.13032e+17 | 2.38754e+18 |
| (756, 'gk_est_log')         | -1.079     | 7.85      | 1.11258e+17 | 2.41615e+18 |
| (756, 'rs_est_log')         | -0.548     | 18.87     | 8.77846e+16 | 1.62212e+18 |
| (1008, 'square_est_log')    | -0.659     | 10.463    | 1.34258e+17 | 2.59699e+18 |
| (1008, 'parkinson_est_log') | -1.024     | 7.822     | 1.37834e+17 | 2.82278e+18 |
| (1008, 'gk_est_log')        | -1.022     | 7.723     | 1.33777e+17 | 2.76729e+18 |
| (1008, 'rs_est_log')        | -0.587     | 16.067    | 1.07053e+17 | 1.98173e+18 |
| (1260, 'square_est_log')    | -0.702     | 10.034    | 1.42917e+17 | 2.54673e+18 |
| (1260, 'parkinson_est_log') | -0.999     | 8.011     | 1.58539e+17 | 3.11109e+18 |
| (1260, 'gk_est_log')        | -1.02      | 7.764     | 1.54594e+17 | 3.04031e+18 |
| (1260, 'rs_est_log')        | -0.666     | 15.929    | 1.22436e+17 | 2.10492e+18 |

**Table 3: Ljung-Box Test Results for HAR Model Residuals**

|                           | square_est_log     | parkinson_est_log | gk_est_log         | rs_est_log         |
|---------------------------|--------------------|-------------------|--------------------|--------------------|
| (252, 'lb_stat_10')       | 16.008475938502087 | 18.09235100877    | 17.70464279279     | 15.392590450654132 |
| (252, 'lb_p_10')          | 0.0993900239809141 | 0.11892929075015  | 0.1533854603460952 | 0.122475306937539  |
| (252, 'lb_stat_20')       | 21.81941500341711  | 16.2844411937372  | 10.06099219339808  | 19.08008865267     |
| (252, 'lb_p_20')          | 0.3503895516427386 | 0.8824076348007   | 0.60585078160539   | 0.9627776490109    |
| (252, 'white_noise_flag') | True               | True              | True               | True               |
| (252, 'lb_lags_used')     | (10, 20)           | (10, 20)          | (10, 20)           | (10, 20)           |
| (252, 'n_obs')            | 4759               | 4759              | 4759               | 4759               |
| (252, 'name')             | square_est_log     | parkinson_est_log | gk_est_log         | rs_est_log         |
| (504, 'lb_stat_10')       | 15.018491131600838 | 15.509887281825   | 10.0670729400835   | 15.2359088932283   |
| (504, 'lb_p_10')          | 0.1313891558874568 | 0.177259582091    | 0.15870962764135   | 0.18171664521069   |
| (504, 'lb_stat_20')       | 20.91016299125434  | 14.46590680508728 | 10.2035636490906   | 17.925612625665    |
| (504, 'lb_p_20')          | 0.4024372413949886 | 0.5549357870577   | 0.6554498464188    | 0.9440931215509    |
| (504, 'white_noise_flag') | True               | True              | True               | True               |

|                            | square_est_log   | parkinson_est_log   | glg_est_log        | rs_est_log      |
|----------------------------|------------------|---------------------|--------------------|-----------------|
| (504, 'lb_lags_used')      | (10, 20)         | (10, 20)            | (10, 20)           | (10, 20)        |
| (504, 'n_obs')             | 4507             | 4507                | 4507               | 4507            |
| (504, 'name')              | square_est_log   | parkinson_est_log   | glg_est_log        | rs_est_log      |
| (756, 'lb_stat_10')        | 20.54020409588   | 10.670852206561782  | 16.8584344157935   | 33288133808632  |
| (756, 'lb_p_10')           | 0.0245379690497  | 1.5696146401087683  | 1.655474050995     | 3634026038526   |
| (756, 'lb_stat_20')        | 24.523774745157  | 12.22199891143532   | 12.2297642374957   | 979330668939435 |
| (756, 'lb_p_20')           | 0.2202608247432  | 1.9082959780239887  | 1.9406238846125    | 357139142221    |
| (756, 'white_noise_flag')  | False            | True                | True               | True            |
| (756, 'lb_lags_used')      | (10, 20)         | (10, 20)            | (10, 20)           | (10, 20)        |
| (756, 'n_obs')             | 4255             | 4255                | 4255               | 4255            |
| (756, 'name')              | square_est_log   | parkinson_est_log   | glg_est_log        | rs_est_log      |
| (1008, 'lb_stat_10')       | 23.2950113332672 | 10.566863022551672  | 10.8687334666162   | 1937061809858   |
| (1008, 'lb_p_10')          | 0.0097084905201  | 1.3209779013640925  | 1.415783980357     | 18895321646346  |
| (1008, 'lb_stat_20')       | 27.727526338080  | 13.74242451199415   | 13.5055835248199   | 643920189448703 |
| (1008, 'lb_p_20')          | 0.1160079273748  | 1.58324488329856    | 1.603965157490348  | 773058591410155 |
| (1008, 'white_noise_flag') | False            | True                | True               | True            |
| (1008, 'lb_lags_used')     | (10, 20)         | (10, 20)            | (10, 20)           | (10, 20)        |
| (1008, 'n_obs')            | 4003             | 4003                | 4003               | 4003            |
| (1008, 'name')             | square_est_log   | parkinson_est_log   | glg_est_log        | rs_est_log      |
| (1260, 'lb_stat_10')       | 19.2443971738432 | 10.5843875226791785 | 11.65049803707     | 172777441706476 |
| (1260, 'lb_p_10')          | 0.0372660519826  | 0.3858176009360295  | 1.0770637800567    | 192548932354614 |
| (1260, 'lb_stat_20')       | 24.352538087781  | 13.872461503536658  | 13.75687945702     | 190072582917874 |
| (1260, 'lb_p_20')          | 0.2273316919990  | 1.79194234380675    | 1.5055488496462855 | 92174784716323  |
| (1260, 'white_noise_flag') | False            | True                | True               | True            |
| (1260, 'lb_lags_used')     | (10, 20)         | (10, 20)            | (10, 20)           | (10, 20)        |
| (1260, 'n_obs')            | 3751             | 3751                | 3751               | 3751            |




Figure 1: ACF and PACF for All Volatility Estimators

(1260, 'name') square\_est\_log parkinson\_est\_log gk\_est\_log rs\_est\_log

---

## HAR Model Predictions vs True RV

The following plots compare the predicted volatility from each estimator against the true realized volatility.

### HAR Model: Predictions vs True RV (Window=252)




Figure 2: HAR Model: Predictions vs True RV (Window=252)

### HAR Model: Predictions vs True RV (Window=504)




Figure 3: HAR Model: Predictions vs True RV (Window=504)

### HAR Model: Predictions vs True RV (Window=756)

### HAR Model: Predictions vs True RV (Window=1008)

### HAR Model: Predictions vs True RV (Window=1260)




Figure 4: HAR Model: Predictions vs True RV (Window=756)




Figure 5: HAR Model: Predictions vs True RV (Window=1008)




Figure 6: HAR Model: Predictions vs True RV (Window=1260)

## Loss Metrics Over Time

QLIKE (Quasi-Likelihood) and MSPE (Mean Squared Prediction Error) are computed over time for each window. These metrics help assess forecast calibration and error magnitude.

### QLIKE Loss Over Time (Window=252)




Figure 7: QLIKE Loss Over Time (Window=252)

### QLIKE Loss Over Time (Window=504)




Figure 8: QLIKE Loss Over Time (Window=504)

### QLIKE Loss Over Time (Window=756)

### QLIKE Loss Over Time (Window=1008)

### QLIKE Loss Over Time (Window=1260)

### MSPE Loss Over Time (Window=252)




Figure 9: QLIKE Loss Over Time (Window=756)




Figure 10: QLIKE Loss Over Time (Window=1008)




Figure 11: QLIKE Loss Over Time (Window=1260)




Figure 12: MSPE Loss Over Time (Window=252)

### MSPE Loss Over Time (Window=504)




Figure 13: MSPE Loss Over Time (Window=504)

### MSPE Loss Over Time (Window=756)

### MSPE Loss Over Time (Window=1008)

### MSPE Loss Over Time (Window=1260)

## Ensemble Model Results

### Ensemble Weights

The ensemble model combines predictions from multiple estimators using inverse QLIKE weighting. Below are the weights assigned to each estimator for different window sizes.






Figure 14: MSPE Loss Over Time (Window=756)




Figure 15: MSPE Loss Over Time (Window=1008)




Figure 16: MSPE Loss Over Time (Window=1260)

**Table 4: Ensemble Model Weights by Window**

| Window | square_est_log | parkinson_est_log | gk_est_log | rs_est_log |
|--------|----------------|-------------------|------------|------------|
| 252    | 0.25           | 0.25              | 0.25       | 0.25       |
| 504    | 0.333          | 0.333             | 0.333      | 0          |
| 756    | 0.25           | 0.25              | 0.25       | 0.25       |
| 1008   | 0.25           | 0.25              | 0.25       | 0.25       |
| 1260   | 0.25           | 0.25              | 0.25       | 0.25       |

**Ensemble Performance Summary****Table 5: Ensemble Model Performance Metrics**

|           | QLIKE_mean | QLIKE_std | MSPE_mean   | MSPE_std    |
|-----------|------------|-----------|-------------|-------------|
| (252, 0)  | -1.296     | 8.109     | 1.73417e+17 | 5.23932e+18 |
| (504, 0)  | -1.26      | 7.615     | 1.21603e+17 | 2.87704e+18 |
| (756, 0)  | -1.113     | 7.987     | 1.05214e+17 | 2.21583e+18 |
| (1008, 0) | -1.051     | 7.824     | 1.25336e+17 | 2.44018e+18 |
| (1260, 0) | -1.048     | 7.819     | 1.41157e+17 | 2.56842e+18 |

**Table 6: Ensemble Model Ljung-Box Test Results (Part 1/2)**

|           | lb_stat_10 | lb_p_10 | lb_stat_20 | lb_p_20 |
|-----------|------------|---------|------------|---------|
| (252, 0)  | 8.273      | 0.602   | 14.126     | 0.824   |
| (252, 1)  | 8.273      | 0.602   | 14.126     | 0.824   |
| (504, 0)  | 6.612      | 0.762   | 12.008     | 0.916   |
| (504, 1)  | 6.612      | 0.762   | 12.008     | 0.916   |
| (756, 0)  | 9.037      | 0.529   | 12.573     | 0.895   |
| (756, 1)  | 9.037      | 0.529   | 12.573     | 0.895   |
| (1008, 0) | 10.321     | 0.413   | 13.816     | 0.84    |
| (1008, 1) | 10.321     | 0.413   | 13.816     | 0.84    |
| (1260, 0) | 10.375     | 0.408   | 14.274     | 0.816   |
| (1260, 1) | 10.375     | 0.408   | 14.274     | 0.816   |

**Table 6: Ensemble Model Ljung-Box Test Results (Part 2/2)**

|          | white_noise_flag | lb_lags_used | n_obs | name |
|----------|------------------|--------------|-------|------|
| (252, 0) | True             | 10           | 4759  |      |
| (252, 1) | True             | 20           | 4759  |      |
| (504, 0) | True             | 10           | 4507  |      |
| (504, 1) | True             | 20           | 4507  |      |

|           | white_noise_flag | lb_lags_used | n_obs | name |
|-----------|------------------|--------------|-------|------|
| (756, 0)  | True             | 10           | 4255  |      |
| (756, 1)  | True             | 20           | 4255  |      |
| (1008, 0) | True             | 10           | 4003  |      |
| (1008, 1) | True             | 20           | 4003  |      |
| (1260, 0) | True             | 10           | 3751  |      |
| (1260, 1) | True             | 20           | 3751  |      |

## Ensemble Predictions and Loss Metrics

### Ensemble Model: Predictions vs True RV (Window=252)




Figure 17: Ensemble Model: Predictions vs True RV (Window=252)

### Ensemble Model: Predictions vs True RV (Window=504)




Figure 18: Ensemble Model: Predictions vs True RV (Window=504)

### Ensemble Model: Predictions vs True RV (Window=756)




Figure 19: Ensemble Model: Predictions vs True RV (Window=756)

### Ensemble Model: Predictions vs True RV (Window=1008)




Figure 20: Ensemble Model: Predictions vs True RV (Window=1008)

### Ensemble Model: Predictions vs True RV (Window=1260)

Ensemble QLIKE Loss (Window=252)

Ensemble QLIKE Loss (Window=504)

Ensemble QLIKE Loss (Window=756)

Ensemble QLIKE Loss (Window=1008)

Ensemble QLIKE Loss (Window=1260)

Ensemble MSPE Loss (Window=252)




Figure 21: Ensemble Model: Predictions vs True RV (Window=1260)




Figure 22: Ensemble QLIKE Loss (Window=252)




Figure 23: Ensemble QLIKE Loss (Window=504)




Figure 24: Ensemble QLIKE Loss (Window=756)




Figure 25: Ensemble QLIKE Loss (Window=1008)




Figure 26: Ensemble QLIKE Loss (Window=1260)




Figure 27: Ensemble MSPE Loss (Window=252)




Figure 28: Ensemble MSPE Loss (Window=504)

### Ensemble MSPE Loss (Window=504)

### Ensemble MSPE Loss (Window=756)




Figure 29: Ensemble MSPE Loss (Window=756)

### Ensemble MSPE Loss (Window=1008)




Figure 30: Ensemble MSPE Loss (Window=1008)

### Ensemble MSPE Loss (Window=1260)

### Diebold-Mariano Test

The Diebold-Mariano (DM) test evaluates whether there is a statistically significant difference in predictive accuracy between two models. Here we compare Windows 504 and 756.

### DM Test Results: Window 504 vs Window 756






Figure 31: Ensemble MSPE Loss (Window=1260)

| Metric       | Value                            |
|--------------|----------------------------------|
| DM Statistic | -1.4633                          |
| P-value      | 0.1434                           |
| Better Model | None (No significant difference) |
| Significant? | 0.0000                           |
| Alpha        | 0.0500                           |
| Observations | 4255.0000                        |

**Interpretation:** With  $p\text{-value} = 0.1434 > 0.05$ , we fail to reject the null hypothesis of equal predictive accuracy. This indicates no statistically significant difference between Window 504 and 756. Either window may be used, and selection can be based on secondary metrics or practical considerations.

## HAR-X Model Results

### Exogenous Variables

The HAR-X model extends the HAR model by incorporating exogenous variables: - **UST10Y**: 10-Year US Treasury Yield - **HYOAS**: High Yield Option-Adjusted Spread - **TermSpread\_10Y\_2Y**: Term Spread (10Y - 2Y) - **VIX**: CBOE Volatility Index - **Breakeven10Y**: 10-Year Breakeven Inflation Rate

All exogenous variables were standardized using expanding window standardization to prevent look-ahead bias.

**Table 7: Diagnostic Tests for Exogenous Variables (After Differencing) (Part 1/3)**

| Estimator         | ADF stat | ADF p | ADF pass (p > ) | KPSS stat |
|-------------------|----------|-------|-----------------|-----------|
| UST10Y            | -3.431   | 0.01  | True            | 1.567     |
| HYOAS             | -3.298   | 0.015 | True            | 0.882     |
| TermSpread_10Y_2Y | -17.061  | 0     | True            | 0.048     |
| VIX               | -5.957   | 0     | True            | 0.463     |
| Breakeven10Y      | -4.528   | 0     | True            | 0.823     |

**Table 7: Diagnostic Tests for Exogenous Variables (After Differencing) (Part 2/3)**

| Estimator         | KPSS p | KPSS pass (p > ) | LB p @10 | LB p @20 |
|-------------------|--------|------------------|----------|----------|
| UST10Y            | 0.01   | False            | 0        | 0        |
| HYOAS             | 0.01   | False            | 0        | 0        |
| TermSpread_10Y_2Y | 0.1    | True             | 0        | 0        |
| VIX               | 0.05   | False            | 0        | 0        |
| Breakeven10Y      | 0.01   | False            | 0        | 0        |

**Table 7: Diagnostic Tests for Exogenous Variables (After Differencing) (Part 3/3)**

| Estimator         | White noise (LB) | Stationary (ADF KPSS) |
|-------------------|------------------|-----------------------|
| UST10Y            | 0                | 0                     |
| HYOAS             | 0                | 0                     |
| TermSpread_10Y_2Y | 0                | 1                     |
| VIX               | 0                | 0                     |
| Breakeven10Y      | 0                | 0                     |

## HARX Model Performance

The HAR-X model performance across different rolling window sizes is presented below.

**Table 8: HAR-X Model Performance Summary**

|                            | QLIKE_mean | QLIKE_std | MSPE_mean   | MSPE_std    |
|----------------------------|------------|-----------|-------------|-------------|
| (252, 'square_est_log')    | -0.831     | 14.471    | 2.9745e+17  | 7.70254e+18 |
| (252, 'parkinson_est_log') | -1.052     | 9.402     | 3.05783e+17 | 8.70434e+18 |
| (252, 'gk_est_log')        | -1.002     | 9.675     | 3.13035e+17 | 9.41853e+18 |
| (252, 'rs_est_log')        | -0.933     | 10.851    | 3.52866e+17 | 1.03554e+19 |
| (504, 'square_est_log')    | -1.249     | 7.798     | 2.24989e+17 | 5.43247e+18 |

|                                | QLIKE_mean | QLIKE_std | MSPE_mean   | MSPE_std    |
|--------------------------------|------------|-----------|-------------|-------------|
| (504,<br>'parkinson_est_log')  | -1.218     | 7.713     | 2.17787e+17 | 5.2861e+18  |
| (504, 'gk_est_log')            | -1.189     | 7.774     | 2.05416e+17 | 4.7914e+18  |
| (504, 'rs_est_log')            | -0.773     | 21.023    | 2.18173e+17 | 5.54433e+18 |
| (756, 'square_est_log')        | -1.308     | 6.913     | 2.49869e+17 | 7.24045e+18 |
| (756,<br>'parkinson_est_log')  | -1.277     | 6.95      | 2.28376e+17 | 6.59822e+18 |
| (756, 'gk_est_log')            | -1.268     | 6.944     | 2.19693e+17 | 6.15746e+18 |
| (756, 'rs_est_log')            | -0.994     | 14.202    | 2.27386e+17 | 6.76839e+18 |
| (1008,<br>'square_est_log')    | -1.181     | 6.962     | 2.34018e+17 | 5.99825e+18 |
| (1008,<br>'parkinson_est_log') | -1.22      | 6.823     | 2.10392e+17 | 5.0361e+18  |
| (1008, 'gk_est_log')           | -1.214     | 6.806     | 2.06946e+17 | 4.99705e+18 |
| (1008, 'rs_est_log')           | -0.993     | 10.734    | 2.08147e+17 | 5.09185e+18 |
| (1260,<br>'square_est_log')    | -1.161     | 7.018     | 2.20667e+17 | 5.07197e+18 |
| (1260,<br>'parkinson_est_log') | -1.201     | 6.922     | 2.12865e+17 | 4.66342e+18 |
| (1260, 'gk_est_log')           | -1.197     | 6.94      | 2.05688e+17 | 4.26511e+18 |
| (1260, 'rs_est_log')           | -1.031     | 9.766     | 1.94173e+17 | 4.05241e+18 |

**Table 9: HAR-X Model Ljung-Box Test Results**

|                              | square_est_log                                               | parkinson_est_log                                            | gk_est_log                                                   | rs_est_log                                                   |
|------------------------------|--------------------------------------------------------------|--------------------------------------------------------------|--------------------------------------------------------------|--------------------------------------------------------------|
| (252,<br>'lb_stat_10')       | 24.70129355780335697236760907940002501925974865999353625651  | 33.5697236760907940002501925974865999353625651               | 33.5697236760907940002501925974865999353625651               | 33.5697236760907940002501925974865999353625651               |
| (252, 'lb_p_10')             | 0.005941386962812366182218980778513951459047096861744676368  | 0.005941386962812366182218980778513951459047096861744676368  | 0.005941386962812366182218980778513951459047096861744676368  | 0.005941386962812366182218980778513951459047096861744676368  |
| (252,<br>'lb_stat_20')       | 31.7453231009012788111156339136512242966143320753549933427   | 31.7453231009012788111156339136512242966143320753549933427   | 31.7453231009012788111156339136512242966143320753549933427   | 31.7453231009012788111156339136512242966143320753549933427   |
| (252, 'lb_p_20')             | 0.0460887785803769462508070964857613626832657087475144834253 | 0.0460887785803769462508070964857613626832657087475144834253 | 0.0460887785803769462508070964857613626832657087475144834253 | 0.0460887785803769462508070964857613626832657087475144834253 |
| (252,<br>'white_noise_flag') | False                                                        | True                                                         | True                                                         | True                                                         |
| (252,<br>'lb_lags_used')     | (10, 20)                                                     | (10, 20)                                                     | (10, 20)                                                     | (10, 20)                                                     |
| (252, 'n_obs')               | 5256                                                         | 5256                                                         | 5256                                                         | 5256                                                         |
| (252, 'name')                | square_est_log                                               | parkinson_est_log                                            | gk_est_log                                                   | rs_est_log                                                   |
| (504,<br>'lb_stat_10')       | 10.7632700735271203628608821662545465859978951725639313884   | 10.7632700735271203628608821662545465859978951725639313884   | 10.7632700735271203628608821662545465859978951725639313884   | 10.7632700735271203628608821662545465859978951725639313884   |
| (504, 'lb_p_10')             | 0.376256512629635060914791949988194818054680018901092376795  | 0.376256512629635060914791949988194818054680018901092376795  | 0.376256512629635060914791949988194818054680018901092376795  | 0.376256512629635060914791949988194818054680018901092376795  |
| (504,<br>'lb_stat_20')       | 17.76771537553255716842695705867575012953218655331486494261  | 17.76771537553255716842695705867575012953218655331486494261  | 17.76771537553255716842695705867575012953218655331486494261  | 17.76771537553255716842695705867575012953218655331486494261  |

|                            | square_est_log     | parkinson_est_log | glg_est_log      | rs_est_log      |
|----------------------------|--------------------|-------------------|------------------|-----------------|
| (504, 'lb_p_20')           | 0.6027068025330451 | 2964094691304791  | 09450569560537   | 66799620039438  |
| (504, 'white_noise_flag')  | True               | True              | True             | True            |
| (504, 'lb_lags_used')      | (10, 20)           | (10, 20)          | (10, 20)         | (10, 20)        |
| (504, 'n_obs')             | 5004               | 5004              | 5004             | 5004            |
| (504, 'name')              | square_est_log     | parkinson_est_log | glg_est_log      | rs_est_log      |
| (756, 'lb_stat_10')        | 16.24155003701079  | 6678241926661056  | 447057293162     | 6861802170436   |
| (756, 'lb_p_10')           | 0.0929233901117383 | 2976183466452364  | 659586082733     | 73479445041     |
| (756, 'lb_stat_20')        | 22.26747874884612  | 673191876647345   | 28423016168.9354 | 7521556253      |
| (756, 'lb_p_20')           | 0.3261746921534065 | 40812676010436    | 4997560940102    | 50236853958826  |
| (756, 'white_noise_flag')  | True               | True              | True             | True            |
| (756, 'lb_lags_used')      | (10, 20)           | (10, 20)          | (10, 20)         | (10, 20)        |
| (756, 'n_obs')             | 4752               | 4752              | 4752             | 4752            |
| (756, 'name')              | square_est_log     | parkinson_est_log | glg_est_log      | rs_est_log      |
| (1008, 'lb_stat_10')       | 20.44766402391143  | 799151954180372   | 2400469190953    | 25791083441896  |
| (1008, 'lb_p_10')          | 0.0252918133164109 | 882343815404053   | 1229195169975    | 425999308923    |
| (1008, 'lb_stat_20')       | 25.876696692775417 | 7979849132908     | 3035575227821592 | 3008163900503   |
| (1008, 'lb_p_20')          | 0.1699229439074254 | 32850366125733    | 1835083850836    | 594386214413116 |
| (1008, 'white_noise_flag') | False              | True              | True             | True            |
| (1008, 'lb_lags_used')     | (10, 20)           | (10, 20)          | (10, 20)         | (10, 20)        |
| (1008, 'n_obs')            | 4500               | 4500              | 4500             | 4500            |
| (1008, 'name')             | square_est_log     | parkinson_est_log | glg_est_log      | rs_est_log      |
| (1260, 'lb_stat_10')       | 18.115190446320057 | 803523667434471   | 13330069836459   | 41638356064     |
| (1260, 'lb_p_10')          | 0.0530513522143520 | 6840697615164003  | 3200457293838    | 4045793821869   |
| (1260, 'lb_stat_20')       | 24.48650383916712  | 778658083024461   | 8559416312817291 | 29809402796     |
| (1260, 'lb_p_20')          | 0.221786275593225  | 267699549650126   | 5355550827191    | 99101705401047  |
| (1260, 'white_noise_flag') | True               | True              | True             | True            |

|                        | square_est_log | parkinson_est_log | gk_est_log | rs_est_log |
|------------------------|----------------|-------------------|------------|------------|
| (1260, 'lb_lags_used') | (10, 20)       | (10, 20)          | (10, 20)   | (10, 20)   |
| (1260, 'n_obs')        | 4248           | 4248              | 4248       | 4248       |
| (1260, 'name')         | square_est_log | parkinson_est_log | gk_est_log | rs_est_log |

### HAR-X Model Predictions vs True RV

The following plots compare the predicted volatility from each estimator with exogenous variables against the true realized volatility.

#### HAR-X Model: Predictions vs True RV (Window=252)




Figure 32: HAR-X Model: Predictions vs True RV (Window=252)

#### HAR-X Model: Predictions vs True RV (Window=504)




Figure 33: HAR-X Model: Predictions vs True RV (Window=504)

### HAR-X Model: Predictions vs True RV (Window=756)




Figure 34: HAR-X Model: Predictions vs True RV (Window=756)

### HAR-X Model: Predictions vs True RV (Window=1008)




Figure 35: HAR-X Model: Predictions vs True RV (Window=1008)

### HAR-X Model: Predictions vs True RV (Window=1260)

#### Loss Metrics Over Time

QLIKE (Quasi-Likelihood) and MSPE (Mean Squared Prediction Error) are computed over time for each window. These metrics help assess forecast calibration and error magnitude for the HAR-X model with exogenous variables.

#### HAR-X QLIKE Loss Over Time (Window=252)

#### HAR-X QLIKE Loss Over Time (Window=504)

#### HAR-X QLIKE Loss Over Time (Window=756)




Figure 36: HAR-X Model: Predictions vs True RV (Window=1260)




Figure 37: HAR-X QLIKE Loss Over Time (Window=252)




Figure 38: HAR-X QLIKE Loss Over Time (Window=504)




Figure 39: HAR-X QLIKE Loss Over Time (Window=756)

#### **HAR-X QLIKE Loss Over Time (Window=1008)**




Figure 40: HAR-X QLIKE Loss Over Time (Window=1008)

#### **HAR-X QLIKE Loss Over Time (Window=1260)**

#### **HAR-X MSPE Loss Over Time (Window=252)**

#### **HAR-X MSPE Loss Over Time (Window=504)**

#### **HAR-X MSPE Loss Over Time (Window=756)**

#### **HAR-X MSPE Loss Over Time (Window=1008)**

#### **HAR-X MSPE Loss Over Time (Window=1260)**






Figure 41: HAR-X QLIKE Loss Over Time (Window=1260)




Figure 42: HAR-X MSPE Loss Over Time (Window=252)




Figure 43: HAR-X MSPE Loss Over Time (Window=504)




Figure 44: HAR-X MSPE Loss Over Time (Window=756)




Figure 45: HAR-X MSPE Loss Over Time (Window=1008)




Figure 46: HAR-X MSPE Loss Over Time (Window=1260)

## HAR-X Ensemble Model Results

### Ensemble Weights

The HAR-X ensemble model combines predictions from multiple estimators using inverse QLIKE weighting. Below are the weights assigned to each estimator for different window sizes.

**Table 10: HAR-X Ensemble Model Weights by Window**

| Window | square_est_log | parkinson_est_log | gk_est_log | rs_est_log |
|--------|----------------|-------------------|------------|------------|
| 252    | 0.25           | 0.25              | 0.25       | 0.25       |
| 504    | 0.25           | 0.25              | 0.25       | 0.25       |
| 756    | 0.25           | 0.25              | 0.25       | 0.25       |
| 1008   | 0.25           | 0.25              | 0.25       | 0.25       |
| 1260   | 0.25           | 0.25              | 0.25       | 0.25       |

### HAR-X Ensemble Performance Summary

**Table 11: HAR-X Ensemble Model Performance Metrics**

|           | QLIKE_mean | QLIKE_std | MSPE_mean   | MSPE_std    |
|-----------|------------|-----------|-------------|-------------|
| (252, 0)  | -1.167     | 9.249     | 2.48679e+17 | 7.04661e+18 |
| (504, 0)  | -1.268     | 7.505     | 1.80769e+17 | 4.25835e+18 |
| (756, 0)  | -1.272     | 6.987     | 1.63743e+17 | 4.44177e+18 |
| (1008, 0) | -1.179     | 6.949     | 1.79972e+17 | 4.61601e+18 |
| (1260, 0) | -1.175     | 7.064     | 1.62983e+17 | 3.19451e+18 |

**Table 12: HAR-X Ensemble Model Ljung-Box Test Results (Part 1/2)**

|           | lb_stat_10 | lb_p_10 | lb_stat_20 | lb_p_20 |
|-----------|------------|---------|------------|---------|
| (252, 0)  | 7.828      | 0.646   | 13.448     | 0.857   |
| (252, 1)  | 7.828      | 0.646   | 13.448     | 0.857   |
| (504, 0)  | 5.97       | 0.818   | 12.122     | 0.912   |
| (504, 1)  | 5.97       | 0.818   | 12.122     | 0.912   |
| (756, 0)  | 9.95       | 0.445   | 14.529     | 0.803   |
| (756, 1)  | 9.95       | 0.445   | 14.529     | 0.803   |
| (1008, 0) | 13.334     | 0.206   | 17.386     | 0.628   |
| (1008, 1) | 13.334     | 0.206   | 17.386     | 0.628   |
| (1260, 0) | 11.796     | 0.299   | 16.221     | 0.703   |
| (1260, 1) | 11.796     | 0.299   | 16.221     | 0.703   |

**Table 12: HAR-X Ensemble Model Ljung-Box Test Results (Part 2/2)**

|           | white_noise_flag | lb_lags_used | n_obs | name |
|-----------|------------------|--------------|-------|------|
| (252, 0)  | True             | 10           | 4756  |      |
| (252, 1)  | True             | 20           | 4756  |      |
| (504, 0)  | True             | 10           | 4504  |      |
| (504, 1)  | True             | 20           | 4504  |      |
| (756, 0)  | True             | 10           | 4252  |      |
| (756, 1)  | True             | 20           | 4252  |      |
| (1008, 0) | True             | 10           | 4000  |      |
| (1008, 1) | True             | 20           | 4000  |      |
| (1260, 0) | True             | 10           | 3748  |      |
| (1260, 1) | True             | 20           | 3748  |      |

#### HAR-X Ensemble Predictions and Loss Metrics

##### HAR-X Ensemble Model: Predictions vs True RV (Window=252)



Figure 47: HAR-X Ensemble Model: Predictions vs True RV (Window=252)

##### HAR-X Ensemble Model: Predictions vs True RV (Window=504)

##### HAR-X Ensemble Model: Predictions vs True RV (Window=756)

##### HAR-X Ensemble Model: Predictions vs True RV (Window=1008)

##### HAR-X Ensemble Model: Predictions vs True RV (Window=1260)

##### HAR-X Ensemble QLIKE Loss (Window=252)

##### HAR-X Ensemble QLIKE Loss (Window=504)

##### HAR-X Ensemble QLIKE Loss (Window=756)




Figure 48: HAR-X Ensemble Model: Predictions vs True RV (Window=504)




Figure 49: HAR-X Ensemble Model: Predictions vs True RV (Window=756)




Figure 50: HAR-X Ensemble Model: Predictions vs True RV (Window=1008)




Figure 51: HAR-X Ensemble Model: Predictions vs True RV (Window=1260)




Figure 52: HAR-X Ensemble QLIKE Loss (Window=252)




Figure 53: HAR-X Ensemble QLIKE Loss (Window=504)




Figure 54: HAR-X Ensemble QLIKE Loss (Window=756)

#### **HAR-X Ensemble QLIKE Loss (Window=1008)**




Figure 55: HAR-X Ensemble QLIKE Loss (Window=1008)

#### **HAR-X Ensemble QLIKE Loss (Window=1260)**

#### **HAR-X Ensemble MSPE Loss (Window=252)**

#### **HAR-X Ensemble MSPE Loss (Window=504)**

#### **HAR-X Ensemble MSPE Loss (Window=756)**

#### **HAR-X Ensemble MSPE Loss (Window=1008)**

#### **HAR-X Ensemble MSPE Loss (Window=1260)**




Figure 56: HAR-X Ensemble QLIKE Loss (Window=1260)




Figure 57: HAR-X Ensemble MSPE Loss (Window=252)




Figure 58: HAR-X Ensemble MSPE Loss (Window=504)






Figure 59: HAR-X Ensemble MSPE Loss (Window=756)




Figure 60: HAR-X Ensemble MSPE Loss (Window=1008)




Figure 61: HAR-X Ensemble MSPE Loss (Window=1260)

**HAR-X Diebold-Mariano Test**

The Diebold-Mariano (DM) test evaluates whether there is a statistically significant difference in predictive accuracy between two HAR-X models. Here we compare Windows 504 and 756.

**HAR-X DM Test Results: Window 504 vs Window 756**

| Metric       | Value                            |
|--------------|----------------------------------|
| DM Statistic | nan                              |
| P-value      | nan                              |
| Better Model | None (No significant difference) |
| Significant? | 0.0000                           |
| Alpha        | 0.0500                           |
| Observations | 4752.0000                        |

**Interpretation:** With p-value = nan, we fail to reject the null hypothesis of equal predictive accuracy. This indicates no statistically significant difference between Window 504 and 756 for the HAR-X model.

**Report generation completed**

*Last Updated: 2025-11-07 at 23:38:30*

**Machine Learning Models**

This section evaluates five state-of-the-art machine learning algorithms for volatility forecasting: - **Random Forest (RF)**: Ensemble of decision trees - **Gradient Boosting Machine (GBM)**: Sequential tree learning with error correction - **XGBoost**: Optimized gradient boosting with GPU support - **LightGBM**: Fast gradient boosting with leaf-wise growth - **CatBoost**: Boosting optimized for categorical features

Each model is trained on a rolling window of historical volatility estimators and exogenous variables, then evaluated on out-of-sample predictions across three window sizes (252, 504, 756 days).

**RF Model Performance**

**Random Forest (RF)** is an ensemble method that builds multiple decision trees and averages their predictions. It’s robust to outliers and naturally captures non-linear relationships in volatility forecasting without requiring feature scaling.

**Table: RF Performance Summary Across Windows (Part 1/2)**

|   | Window | QLIKE | QLIKE Std | MSPE    |
|---|--------|-------|-----------|---------|
| 0 | full   | 7.522 | 29.592    | 4391.26 |

**Table: RF Performance Summary Across Windows (Part 2/2)**

|   | MSPE Std | RMSE  | Dir. Accuracy |
|---|----------|-------|---------------|
| 0 | 55665    | 0.045 | 51.58%        |

**RF Predictions vs True RV** The following plots compare predicted volatility against true realized volatility for each rolling window.

#### RF: Predictions vs True RV (Window=full)




Figure 62: RF: Predictions vs True RV (Window=full)

**RF Loss Metrics Over Time** QLIKE (Quasi-Likelihood) and MSPE (Mean Squared Prediction Error) are computed over time for each window. These metrics help assess forecast calibration and error magnitude for the ML model across different training periods.

#### RF QLIKE Loss Over Time (Window=full)

#### RF MSPE Loss Over Time (Window=full)

#### GBM Model Performance

**Gradient Boosting Machine (GBM)** sequentially builds decision trees, with each tree correcting the errors of previous ones. This incremental approach often provides superior performance for forecasting tasks by explicitly learning residual patterns.




Figure 63: RF QLIKE Loss Over Time (Window=full)




Figure 64: RF MSPE Loss Over Time (Window=full)

**Table: GBM Performance Summary Across Windows (Part 1/2)**

|   | Window | QLIKE  | QLIKE Std | MSPE    |
|---|--------|--------|-----------|---------|
| 0 | full   | 24.459 | 146.674   | 9738.12 |

**Table: GBM Performance Summary Across Windows (Part 2/2)**

|   | MSPE Std | RMSE  | Dir. Accuracy |
|---|----------|-------|---------------|
| 0 | 159979   | 0.045 | 50.53%        |

**GBM Predictions vs True RV** The following plots compare predicted volatility against true realized volatility for each rolling window.

#### GBM: Predictions vs True RV (Window=full)




Figure 65: GBM: Predictions vs True RV (Window=full)

**GBM Loss Metrics Over Time** QLIKE (Quasi-Likelihood) and MSPE (Mean Squared Prediction Error) are computed over time for each window. These metrics help assess forecast calibration and error magnitude for the ML model across different training periods.

#### GBM QLIKE Loss Over Time (Window=full)

#### GBM MSPE Loss Over Time (Window=full)

#### XGBOOST Model Performance

**XGBoost** (eXtreme Gradient Boosting) is an optimized gradient boosting implementation with GPU acceleration support. It includes regularization to prevent overfitting and has proven highly effective for time series forecasting competitions.




Figure 66: GBM QLIKE Loss Over Time (Window=full)




Figure 67: GBM MSPE Loss Over Time (Window=full)

Table: XGBOOST Performance Summary Across Windows (Part 1/2)

|   | Window | QLIKE  | QLIKE Std | MSPE    |
|---|--------|--------|-----------|---------|
| 0 | full   | 14.502 | 64.309    | 7132.32 |

Table: XGBOOST Performance Summary Across Windows (Part 2/2)

|   | MSPE Std | RMSE  | Dir. Accuracy |
|---|----------|-------|---------------|
| 0 | 135627   | 0.046 | 45.05%        |

**XGBOOST Predictions vs True RV** The following plots compare predicted volatility against true realized volatility for each rolling window.

**XGBOOST: Predictions vs True RV (Window=full)**




Figure 68: XGBOOST: Predictions vs True RV (Window=full)

**XGBOOST Loss Metrics Over Time** QLIKE (Quasi-Likelihood) and MSPE (Mean Squared Prediction Error) are computed over time for each window. These metrics help assess forecast calibration and error magnitude for the ML model across different training periods.

**XGBOOST QLIKE Loss Over Time (Window=full)**

**XGBOOST MSPE Loss Over Time (Window=full)**




Figure 69: XGBOOST QLIKE Loss Over Time (Window=full)




Figure 70: XGBOOST MSPE Loss Over Time (Window=full)



## LIGHTGBM Model Performance

**LightGBM** (Light Gradient Boosting Machine) is a fast, distributed gradient boosting framework with GPU support. It uses leaf-wise tree growth for better accuracy and is especially efficient with large feature sets.

**Table: LIGHTGBM Performance Summary Across Windows (Part 1/2)**

|   | Window | QLIKE | QLIKE Std | MSPE    |
|---|--------|-------|-----------|---------|
| 0 | full   | 9.429 | 31.96     | 4642.25 |

**Table: LIGHTGBM Performance Summary Across Windows (Part 2/2)**

|   | MSPE Std | RMSE  | Dir. Accuracy |
|---|----------|-------|---------------|
| 0 | 66008.1  | 0.045 | 49.68%        |

**LIGHTGBM Predictions vs True RV** The following plots compare predicted volatility against true realized volatility for each rolling window.

### LIGHTGBM: Predictions vs True RV (Window=full)




Figure 71: LIGHTGBM: Predictions vs True RV (Window=full)

**LIGHTGBM Loss Metrics Over Time** QLIKE (Quasi-Likelihood) and MSPE (Mean Squared Prediction Error) are computed over time for each window. These metrics help assess forecast calibration and error magnitude for the ML model across different training periods.

### LIGHTGBM QLIKE Loss Over Time (Window=full)




Figure 72: LIGHTGBM QLIKE Loss Over Time (Window=full)

### LIGHTGBM MSPE Loss Over Time (Window=full)




Figure 73: LIGHTGBM MSPE Loss Over Time (Window=full)

### CATBOOST Model Performance

**CatBoost** (Categorical Boosting) is designed to handle categorical features effectively with GPU acceleration. It uses ordered boosting and symmetric trees to reduce overfitting, making it ideal for heterogeneous financial data.

**Table: CATBOOST Performance Summary Across Windows (Part 1/2)**

|   | Window | QLIKE | QLIKE Std | MSPE    |
|---|--------|-------|-----------|---------|
| 0 | full   | 2.108 | 15.197    | 18453.7 |

**Table: CATBOOST Performance Summary Across Windows (Part 2/2)**

|   | MSPE Std | RMSE  | Dir. Accuracy |
|---|----------|-------|---------------|
| 0 | 312392   | 0.045 | 50.53%        |

**CATBOOST Predictions vs True RV** The following plots compare predicted volatility against true realized volatility for each rolling window.

#### CATBOOST: Predictions vs True RV (Window=full)




Figure 74: CATBOOST: Predictions vs True RV (Window=full)

**CATBOOST Loss Metrics Over Time** QLIKE (Quasi-Likelihood) and MSPE (Mean Squared Prediction Error) are computed over time for each window. These metrics help assess forecast calibration and error magnitude for the ML model across different training periods.

#### CATBOOST QLIKE Loss Over Time (Window=full)

#### CATBOOST MSPE Loss Over Time (Window=full)

---

### ML Models Analysis Completed Successfully

*Last Updated: 2025-11-07 at 23:38:47*

### Temporal Fusion Transformer (TFT) Results

The Temporal Fusion Transformer is a state-of-the-art deep learning architecture for multi-horizon time series forecasting. It combines:




Figure 75: CATBOOST QLIKE Loss Over Time (Window=full)




Figure 76: CATBOOST MSPE Loss Over Time (Window=full)

- **Multi-head attention mechanism:** Captures complex temporal dependencies
- **Variable selection networks:** Automatic feature importance learning
- **Gated residual networks:** Non-linear processing with skip connections
- **Quantile forecasting:** Provides prediction intervals (10th, 50th, 90th percentiles)

**Key Improvements Implemented:** - **Multiple quantiles:** Generates prediction intervals for uncertainty quantification - **Increased model capacity:** Larger hidden sizes and attention heads for better learning - **Enhanced regularization:** Higher dropout to prevent overfitting - **Extended lookback:** 90-day encoder length for capturing longer-term patterns

**TFT Architecture Details:** - Hidden size: 64 (optimized for dataset size) - Attention heads: 4 - Encoder length: 90 days (quarterly lookback) - Dropout: 0.1 (enhanced regularization) - Quantiles: 0.05, 0.1, 0.25, 0.5, 0.75, 0.9, 0.95 (7 quantiles) - Early stopping: Patience of 20 epochs

#### TFT Model Performance (Validation Set)

| Metric               | Value            |
|----------------------|------------------|
| QLIKE Mean           | -2.533680        |
| QLIKE Std            | 1.219096         |
| QLIKE Min            | -5.174289        |
| QLIKE Max            | 7.203637         |
| MSPE Mean            | 1569516.258582   |
| MSPE Std             | 24253059.812042  |
| MSPE Min             | 0.000001         |
| MSPE Max             | 526681836.794077 |
| RMSE                 | 0.047893         |
| Directional Accuracy | 0.4406 (44.06%)  |
| Training Samples     | 4523             |
| Validation Samples   | 530              |

**Figure: TFT Main Performance Analysis (Validation Set)**

**Figure: TFT QLIKE Distribution Analysis**

**Figure: TFT MSPE Distribution Analysis**

**Figure: TFT RMSE Analysis**

---

**Report generation completed**

*Last Updated: 2025-11-07 at 23:40:20*



Figure 77: Figure: TFT Main Performance Analysis (Validation Set)




Figure 78: Figure: TFT QLIKE Distribution Analysis




Figure 79: Figure: TFT MSPE Distribution Analysis




Figure 80: Figure: TFT RMSE Analysis